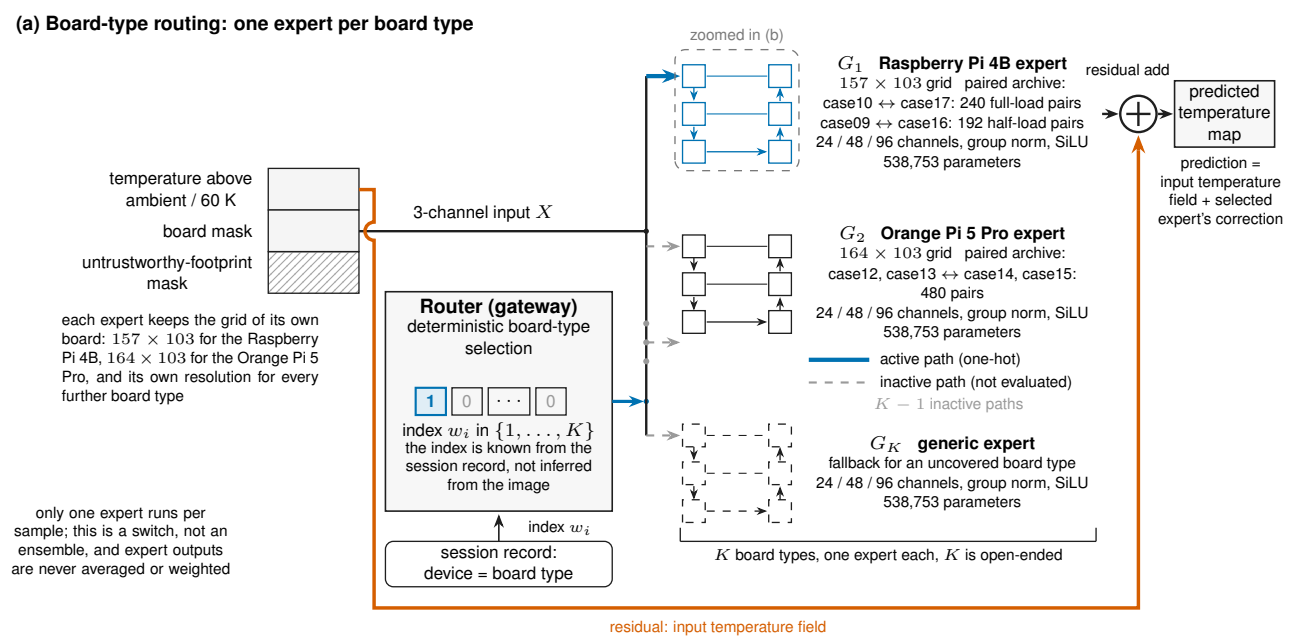

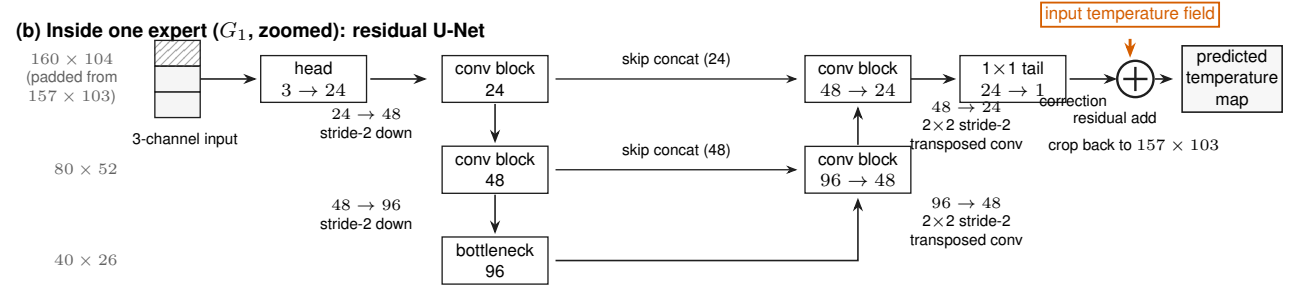

conv block: two 3×3 convolutions with group norm and SiLU

level 2 runs its conv block at 96 channels, then the bottleneck block (also 96); the 1×1 tail is linear

identical topology, independent weights; every expert keeps its own board grid

loss: complete-map L1 over board pixels, in kelvin

input reflect-padded from 157×103 to 160×104 (a multiple of 8); prediction cropped back to 157×103